



Detailed Report


Model Config: gemma-3-1b-it_tensorrt_llm_bls_config_79



Latency Breakdown for Online Performance of gemma-3-1b-it_tensorrt_llm_bls_config_79




GPU Memory vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_79



GPU Utilization vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_79

GPU Power vs. Latency



GPU Power vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_79

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
60	10853.342	0.0	1778.461	0.064	6605.271	7.96271	22106.079232	95.0
64	10837.816	0.0	2188.905	0.057	6572.372	7.96102	22106.079232	95.2
58	9879.954	0.0	1650.611	0.058	6531.953	7.96003	22106.079232	92.4
56	9874.145	0.0	1441.319	0.059	6526.592	7.96223	22106.079232	95.1
59	9863.918	0.0	1702.12	0.07	6510.063	7.96226	22106.079232	95.2
48	9226.814	0.0	724.838	0.081	6358.586	7.74249	22106.079232	94.3
32	7872.28	0.0	147.68	0.048	3932.65	7.95768	22106.079232	95.0
16	3931.181	0.0	0.037	0.047	1971.365	7.96336	22106.079232	95.1
8	1004.519	0.0	0.032	0.045	995.276	7.96254	22106.079232	92.4
1	989.227	0.0	0.044	0.04	877.624	1.1383	22106.079232	92.4
4	915.509	0.0	0.03	0.046	900.751	4.43467	22106.079232	92.5
2	891.034	0.0	0.032	0.047	884.148	2.21982	22106.079232	92.6

The model config **gemma-3-1b-it_tensorrt_llm_bls_config_79** uses 28 GPU instances with a max batch size of 2 and has dynamic batching enabled. 12 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the

measurements taken for this model config in decreasing order of latency.