



Detailed Report


Model Config: gemma-3-1b-it_tensorrt_llm_bls_config_82



Latency Breakdown for Online Performance of gemma-3-1b-it_tensorrt_llm_bls_config_82




GPU Memory vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_82



GPU Utilization vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_82

GPU Power vs. Latency



GPU Power vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_82

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
61	10872.053	0.0	1793.01	0.061	6635.486	7.96169	22106.079232	95.1
62	10858.382	0.0	1836.761	0.067	6709.983	7.96339	22106.079232	92.1
64	10830.63	0.0	2143.424	0.06	6700.637	7.96222	22106.079232	95.3
60	9871.97	0.0	1655.135	0.063	6701.308	7.96285	22106.079232	92.2
56	9868.955	0.0	1269.918	0.061	6717.211	7.96275	22106.079232	95.2
48	9237.289	0.0	735.367	0.058	6184.888	7.74286	22106.079232	92.4
32	6896.592	0.0	75.428	0.048	3998.492	7.96102	22106.079232	95.2
16	2980.775	0.0	0.039	0.047	1966.226	7.96243	22106.079232	95.2
8	1003.969	0.0	0.034	0.048	993.805	7.96292	22106.079232	92.2
1	999.481	0.0	0.042	0.041	876.977	1.1383	22106.079232	92.8
4	916.337	0.0	0.03	0.046	900.617	4.43461	22106.079232	92.6
2	890.251	0.0	0.031	0.044	884.27	2.21988	22106.079232	92.6

The model config **`gemma-3-1b-it_tensorrt_llm_bls_config_82`** uses 29 GPU instances with a max batch size of 2 and has dynamic batching enabled. 12 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the

measurements taken for this model config in decreasing order of latency.