



Detailed Report


Model Config: gemma-3-1b-it_tensorrt_llm_bls_config_106




Latency Breakdown for Online Performance of gemma-3-1b-it_tensorrt_llm_bls_config_106



GPU Memory vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_106



GPU Utilization vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_106



GPU Power vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_106

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
64	7905.22	0.0	2576.603	0.041	4552.383	7.96014	22108.176384	92.2
32	3966.631	0.0	0.088	0.049	3790.233	7.96173	22108.176384	94.8
16	1991.798	0.0	0.039	0.048	1950.476	7.96277	22108.176384	94.9
8	1005.008	0.0	0.037	0.048	997.73	7.96142	22108.176384	94.9
1	989.519	0.0	0.048	0.039	879.751	1.11051	22108.176384	94.8
4	918.737	0.0	0.03	0.047	903.218	4.32348	22108.176384	95.1
2	893.814	0.0	0.032	0.046	887.389	2.21985	22108.176384	95.2

The model config **gemma-3-1b-it_tensorrt_llm_bls_config_106** uses 39 GPU instances with a max batch size of 1 and has dynamic batching enabled. 7 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.