

Online Result Summary

Model: gemma-3-1b-it_tensorrt_llm_bls

GPU(s): 1 x NVIDIA GeForce RTX 4090


Total Available GPU Memory: 23.5 GB

Constraint targets: Max p99 Latency: 10000 ms, Max p99 Time To First Token: 10000 ms


In 1069 measurements across 155 configurations, **gemma-3-1b-it_tensorrt_llm_bls_config_82** is **308%** better than the default configuration at meeting the objectives, under the given constraints, on GPU(s) 1 x NVIDIA GeForce RTX 4090.

- gemma-3-1b-it_tensorrt_llm_bls_config_82**: 29 GPU instances with a max batch size of 2 on platform python

Curves corresponding to the 3 best model configuration(s) out of a total of 155 are shown in the plots.



Throughput vs. Latency curves for 3 best configurations.



GPU Memory vs. Latency curves for 3 best configurations.

The following table summarizes each configuration at the measurement that optimizes the desired metrics under the given constraints.

Model Config Name	Max Batch Size	Dynamic Batching	Total Instance Count	p99 Latency (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
gemma-3-1b-it_tensorrt_llm_bls_config_82	2	Enabled	29:GPU	1003.969	7.96292	22106	92.2
gemma-3-1b-it_tensorrt_llm_bls_config_79	2	Enabled	28:GPU	1004.519	7.96254	22106	92.4
gemma-3-1b-it_tensorrt_llm_bls_config_106	1	Enabled	39:GPU	1005.008	7.96142	22108	94.9
gemma-3-1b-it_tensorrt_llm_bls_config_default	8	Disabled	1:CPU	990.196	1.11048	22101	92.5