



Detailed Report


Model Config: gemma-3-1b-it_tensorrt_llm_bls_config_0



Latency Breakdown for Online Performance of gemma-3-1b-it_tensorrt_llm_bls_config_0




GPU Memory vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_0



GPU Utilization vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_0

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
8	6966.911	0.0	4226.319	0.08	3419.434	1.13828	24220.008448	92.5
4	3488.735	0.0	1633.389	0.065	2569.89	1.13829	24220.008448	95.0
2	1747.741	0.0	849.562	0.033	871.589	1.13828	24220.008448	95.0
1	996.828	0.0	0.029	0.036	879.119	1.11036	24220.008448	92.3

GPU Power vs. Latency



GPU Power vs. Latency curves for config `gemma-3-1b-it_tensorrt_llm_bls_config_0`

The model config **`gemma-3-1b-it_tensorrt_llm_bls_config_0`** uses 1 GPU instance with a max batch size of 4 and has dynamic batching enabled. 4 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.