Detailed Report

Model Config: gemma-3-1b-it_tensorrt_llm_bls_config_default

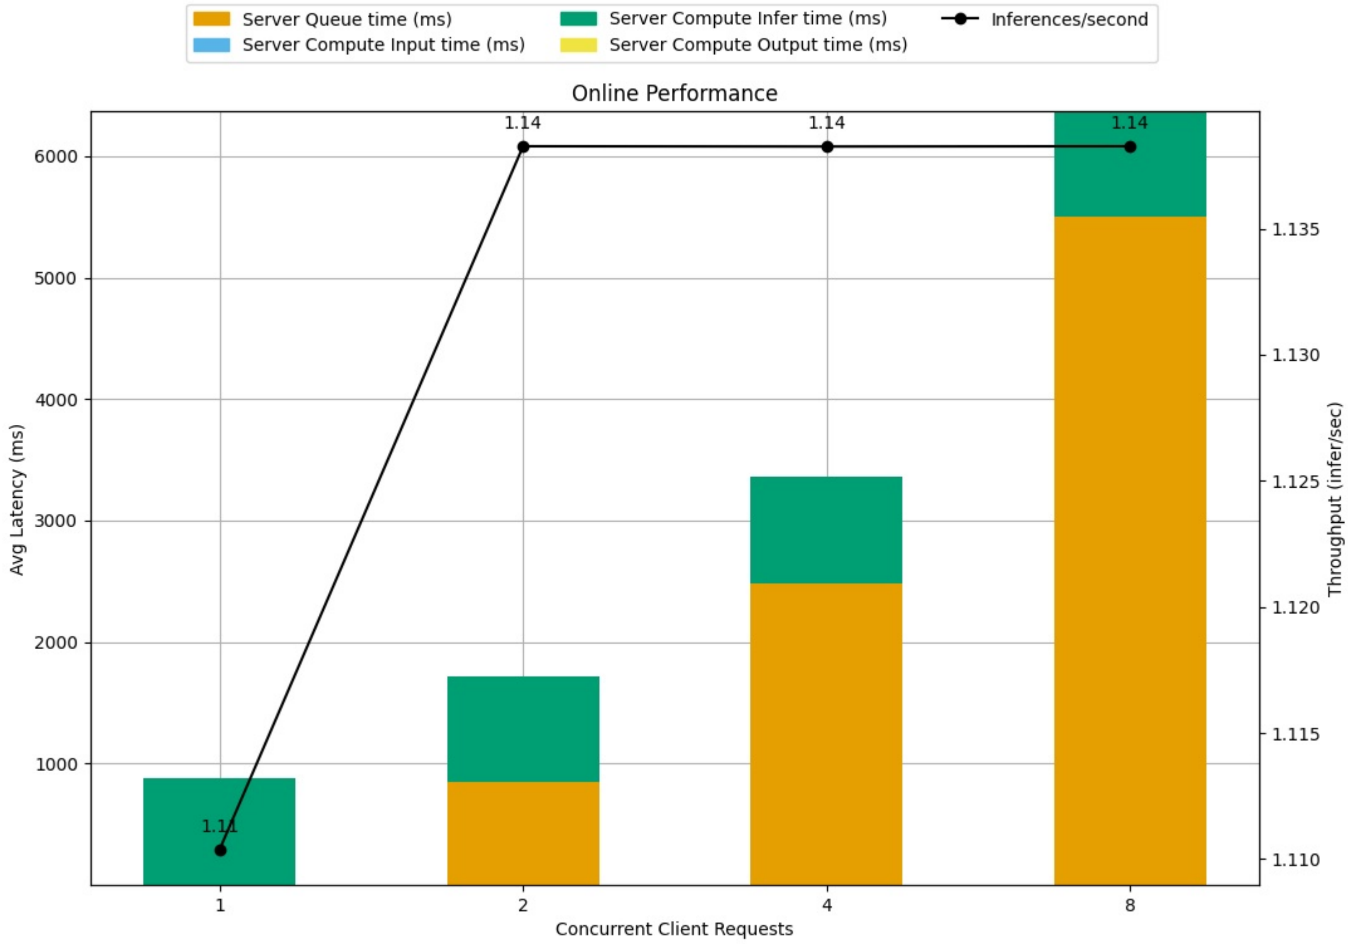

Latency Breakdown for Online Performance of gemma-3-1b-it_tensorrt_llm_bls_config_default

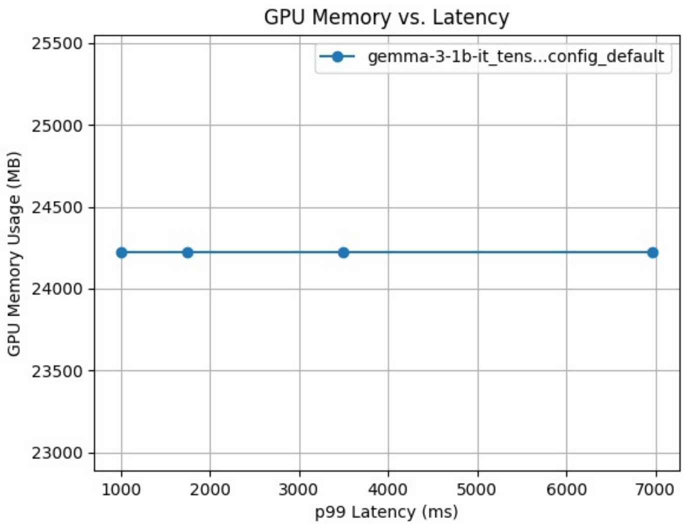

GPU Memory vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_default

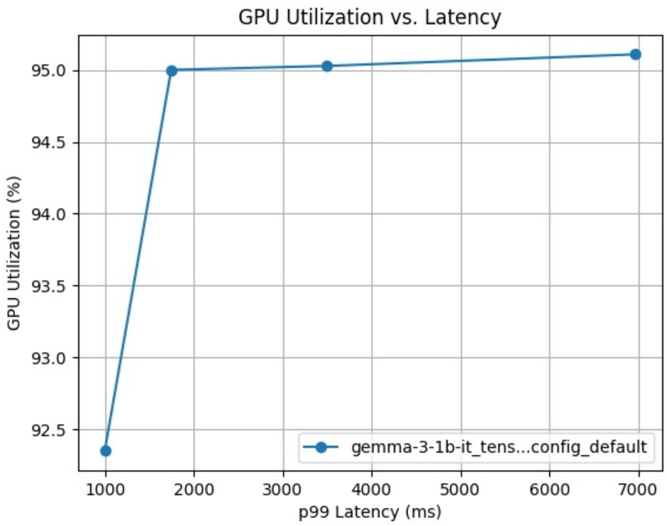

GPU Utilization vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_default

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
8	6971.201	0.0	5501.387	0.034	870.867	1.13827	24220.008448	95.1
4	3490.085	0.0	2486.884	0.036	871.521	1.13826	24220.008448	95.0
2	1746.228	0.0	849.029	0.032	871.029	1.13827	24220.008448	95.0
1	998.472	0.0	0.016	0.036	879.063	1.11037	24220.008448	92.4

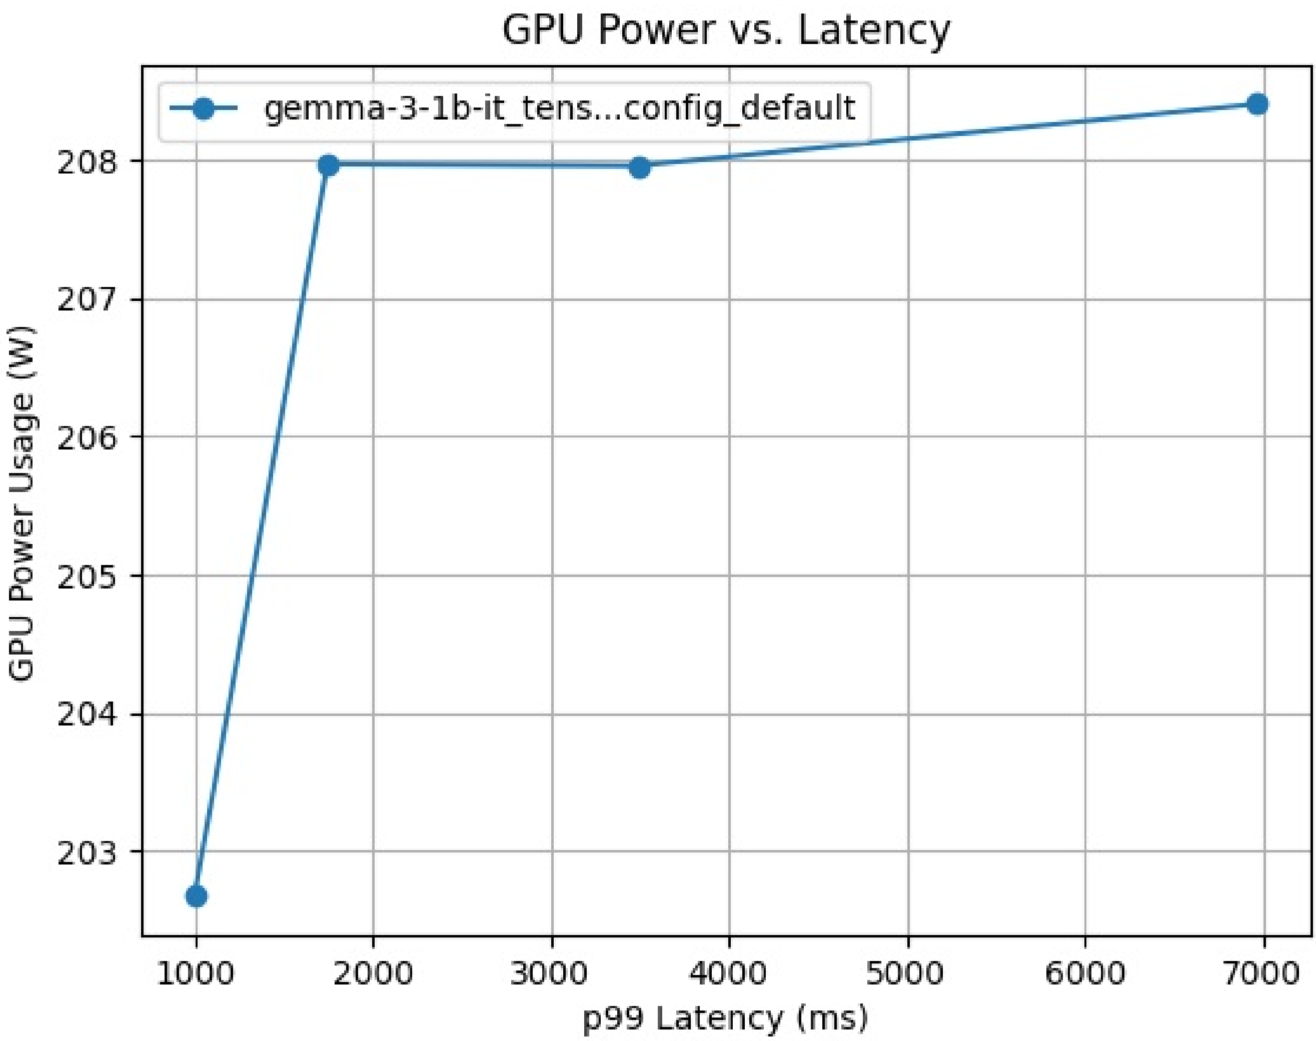

GPU Power vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_default

The model config **gemma-3-1b-it_tensorrt_llm_bls_config_default** uses 1 CPU instance with a max batch size of 8 and has dynamic batching disabled. 4 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.