

Online Result Summary

Model: gemma-3-1b-it_tensorrt_llm_bls

GPU(s): 1 x NVIDIA GeForce RTX 4090


Total Available GPU Memory: 23.5 GB

Constraint targets: Max p99 Latency: 10000 ms, Max p99 Time To First Token: 10000 ms


In 8 measurements across 2 configurations, **gemma-3-1b-it_tensorrt_llm_bls_config_0** provides no gain over the default configuration, under the given constraints, on GPU(s) 1 x NVIDIA GeForce RTX 4090.

- gemma-3-1b-it_tensorrt_llm_bls_config_0:** 1 GPU instance with a max batch size of 4 on platform python

Curves corresponding to the 2 best model configuration(s) out of a total of 2 are shown in the plots.



Throughput vs. Latency curves for 2 best configurations.



GPU Memory vs. Latency curves for 2 best configurations.

The following table summarizes each configuration at the measurement that optimizes the desired metrics under the given constraints.

Model Config Name	Max Batch Size	Dynamic Batching	Total Instance Count	p99 Latency (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
gemma-3-1b-it_tensortt_llm_bls_config_0	4	Enabled	1:GPU	996.828	1.11036	24220	92.3
gemma-3-1b-it_tensortt_llm_bls_config_default	8	Disabled	1:CPU	998.472	1.11037	24220	92.4