



Detailed Report


Model Config: gemma-3-1b-it_tensorrt_llm_bls_config_10



Latency Breakdown for Online Performance of gemma-3-1b-it_tensorrt_llm_bls_config_10




GPU Memory vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_10



GPU Utilization vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_10

GPU Power vs. Latency

GPU Power vs. Latency curves for config **gemma-3-1b-it_tensorrt_llm_bls_config_10**

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
256	35887.05	0.0	619.126	0.2	9506.662	7.99923	22103.98208	96.1
1	998.356	0.0	0.044	0.044	875.809	1.13883	22103.98208	92.7
8	991.91	0.0	0.024	0.034	980.917	7.99926	22103.98208	96.2
4	911.636	0.0	0.031	0.044	893.891	4.44412	22103.98208	92.9
2	887.65	0.0	0.034	0.044	880.751	2.22208	22103.98208	93.1

The model config **gemma-3-1b-it_tensorrt_llm_bls_config_10** uses 8 GPU instances with a max batch size of 64 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.