



# Detailed Report


Model Config: gemma-3-1b-it\_tensorrt\_llm\_bls\_config\_11



Latency Breakdown for Online Performance of gemma-3-1b-it\_tensorrt\_llm\_bls\_config\_11




GPU Memory vs. Latency curves for config gemma-3-1b-it\_tensorrt\_llm\_bls\_config\_11



GPU Utilization vs. Latency curves for config gemma-3-1b-it\_tensorrt\_llm\_bls\_config\_11

GPU Power vs. Latency

GPU Power vs. Latency curves for config `gemma-3-1b-it_tensorrt_llm_bls_config_11`

| Request Concurrency | p99 Latency (ms) | Client Response Wait (ms) | Server Queue (ms) | Server Compute Input (ms) | Server Compute Infer (ms) | Throughput (infer/sec) | Max GPU Memory Usage (MB) | Average GPU Utilization (%) |
|---------------------|------------------|---------------------------|-------------------|---------------------------|---------------------------|------------------------|---------------------------|-----------------------------|
| 256                 | 34893.202        | 0.0                       | 368.022           | 0.174                     | 9285.558                  | 7.99897                | 22103.98208               | 93.4                        |
| 8                   | 993.287          | 0.0                       | 0.027             | 0.037                     | 982.41                    | 7.99922                | 22103.98208               | 93.1                        |
| 1                   | 993.237          | 0.0                       | 0.044             | 0.048                     | 875.525                   | 1.13883                | 22103.98208               | 95.2                        |
| 4                   | 908.883          | 0.0                       | 0.031             | 0.045                     | 892.2                     | 4.4441                 | 22103.98208               | 93.2                        |
| 2                   | 887.299          | 0.0                       | 0.04              | 0.049                     | 880.244                   | 2.22207                | 22103.98208               | 93.0                        |

The model config **`gemma-3-1b-it_tensorrt_llm_bls_config_11`** uses 9 GPU instances with a max batch size of 128 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.