



Detailed Report


Model Config: gemma-3-1b-it_tensorrt_llm_bls_config_12



Latency Breakdown for Online Performance of gemma-3-1b-it_tensorrt_llm_bls_config_12




GPU Memory vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_12



GPU Utilization vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_12

GPU Power vs. Latency



GPU Power vs. Latency curves for config gemma-3-1b-it_tensorrt_llm_bls_config_12

Request Concurrency	p99 Latency (ms)	Client Response Wait (ms)	Server Queue (ms)	Server Compute Input (ms)	Server Compute Infer (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
256	33807.38	0.0	332.912	0.157	10392.258	7.99929	22103.98208	96.1
8	991.83	0.0	0.029	0.04	981.563	7.99921	22103.98208	96.0
4	897.678	0.0	0.033	0.046	891.607	4.44403	22103.98208	95.9
2	883.24	0.0	0.033	0.049	877.234	2.27762	22103.98208	92.9
1	873.486	0.0	0.042	0.054	868.688	1.13882	22103.98208	95.5

The model config **`gemma-3-1b-it_tensorrt_llm_bls_config_12`** uses 10 GPU instances with a max batch size of 128 and has dynamic batching enabled. 5 measurement(s) were obtained for the model config on GPU(s) 1 x NVIDIA GeForce RTX 4090 with total memory 23.5 GB. This model uses the platform .

The first plot above shows the breakdown of the latencies in the latency throughput curve for this model config. Following that are the requested configurable plots showing the relationship between various metrics measured by the Model Analyzer. The above table contains detailed data for each of the measurements taken for this model config in decreasing order of latency.