

Online Result Summary

Model: gemma-3-1b-it_tensorrt_llm_bls

GPU(s): 1 x NVIDIA GeForce RTX 4090


Total Available GPU Memory: 23.5 GB

Constraint targets: None


In 29 measurements across 13 configurations, **gemma-3-1b-it_tensorrt_llm_bls_config_12** is **602%** better than the default configuration at maximizing throughput, under the given constraints, on GPU(s) 1 x NVIDIA GeForce RTX 4090.

- gemma-3-1b-it_tensorrt_llm_bls_config_12**: 10 GPU instances with a max batch size of 128 on platform python

Curves corresponding to the 3 best model configuration(s) out of a total of 13 are shown in the plots.



Throughput vs. Latency curves for 3 best configurations.



GPU Memory vs. Latency curves for 3 best configurations.

The following table summarizes each configuration at the measurement that optimizes the desired metrics under the given constraints.

Model Config Name	Max Batch Size	Dynamic Batching	Total Instance Count	p99 Latency (ms)	Throughput (infer/sec)	Max GPU Memory Usage (MB)	Average GPU Utilization (%)
gemma-3-1b-it_tensorrt_llm_bls_config_12	128	Enabled	10:GPU	33807.38	7.99929	22103	96.1
gemma-3-1b-it_tensorrt_llm_bls_config_10	64	Enabled	8:GPU	991.91	7.99926	22103	96.2
gemma-3-1b-it_tensorrt_llm_bls_config_11	128	Enabled	9:GPU	993.287	7.99922	22103	93.1
gemma-3-1b-it_tensorrt_llm_bls_config_default	8	Disabled	1:CPU	988.39	1.13884	22101	95.8